




Evaluating the CO2 Emission from Gasoline- Powered Light- Duty Vehicles

Group 5 - Applied Machine Learning

Han Chau - 20012654

Randy Duong - 20019147

Nihar Dugade - 20018679





Content

- Problem and research question
- Data description and insights
- Data pre-processing
- Method
- Analysis and results
- Conclusion



Problem and research question

Can we accurately predict whether the CO₂ emissions of a car exceed the allowed threshold using machine learning models?









Can the chosen model be used as a practical tool for monitoring and controlling CO₂ emissions in real-time scenarios?

Data description and insights


- Taken from Kaggle with 22,556 examples and 13 columns.
- Datasets provide model-specific fuel consumption ratings and estimated carbon dioxide emissions for new light-duty vehicles for retail sale in Canada.

	YEAR	MAKE	MODEL	VEHICLE CLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUEL	FUEL CONSUMPTION	HWY (L/100 km)	COMB (L/100 km)	COMB (mpg)	EMISSIONS
0	2000	ACURA	1.6EL	COMPACT	1.6	4	A4	X	9.2	6.7	8.1	35	186
1	2000	ACURA	1.6EL	COMPACT	1.6	4	M5	X	8.5	6.5	7.6	37	175
2	2000	ACURA	3.2TL	MID-SIZE	3.2	6	AS5	Z	12.2	7.4	10.0	28	230
3	2000	ACURA	3.5RL	MID-SIZE	3.5	6	A4	Z	13.4	9.2	11.5	25	264
4	2000	ACURA	INTEGRA	SUBCOMPACT	1.8	4	A4	X	10.0	7.0	8.6	33	198

Data description and insights



Data description and insights



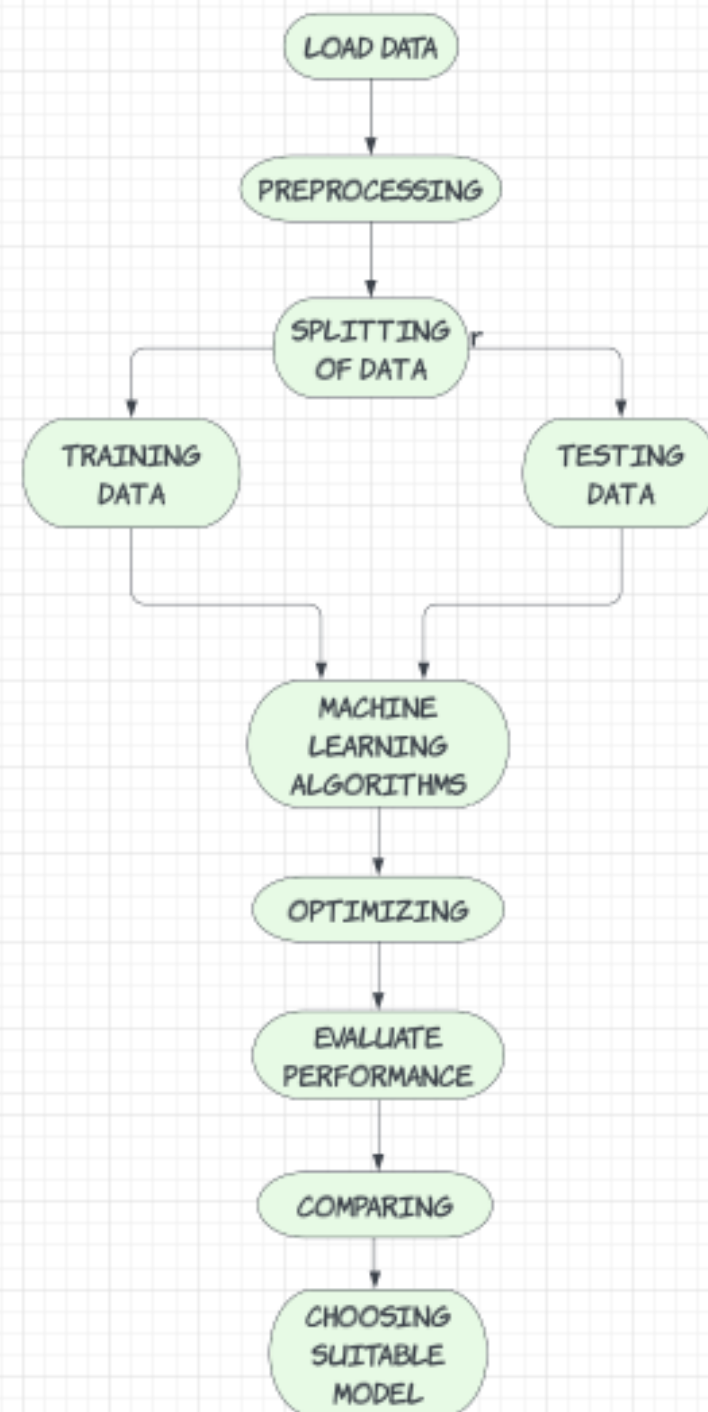
Data pre-processing

- Drop examples which the types of fuel is not Gasoline.
- Transform the emission columns into binary values (1 and 0) depending on whether they exceed the threshold.

$$CO_2 \text{ emissions per km} = \frac{CO_2 \text{ per gallon}}{MPG} = \frac{8,887}{35.4} = 251 \text{ grams}$$

- Eliminate outliers by considering the interquartile range (IQR) of the fuel consumption variable.


METHOD





Machine Learning Algorithms

- Logistic Regression
- Linear SVC
- Random Forest
- Neural Network
- Naïve Bayes



Logistic regression model

Set default for the initial parameters.

The Accuracy is : 97.44 %

Classification report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1217
1	0.97	0.97	0.97	856
accuracy			0.97	2073
macro avg	0.97	0.97	0.97	2073
weighted avg	0.97	0.97	0.97	2073

Logistic regression model

Optimizing the model with GridSearchCV

```
lr_model = LogisticRegression()

# Define hyperparameters to search over
hyperparameters = {
    'penalty': ['l1', 'l2', None],
    'C': [0.01, 0.1, 1, 10, 100, 1000]
}

# Use GridSearchCV to find best hyperparameters
grid_search = GridSearchCV(lr_model, hyperparameters, cv=5)
grid_search.fit(train_x, train_y)
```

Best hyperparameters: {'C': 0.01, 'penalty': None}

Best accuracy score: 0.9868131868131869

Logistic regression model

Optimizing the model with SelectKBest and GridSearchCV

```
pipeline = Pipeline([
    ('select', SelectKBest(score_func=f_classif)), # Select top features using ANOVA F-value
    ('scale', StandardScaler()), # Standardize the data
    ('classify', LogisticRegression()) # Classifier
])
para = {
    'select__k': [5, 8, 10],
    'classify__penalty': ['l1', 'l2', None],
    'classify__C': [0.01, 0.1, 1, 10, 100, 1000]
}

grid_search = GridSearchCV(pipeline, para, cv=5)
grid_search.fit(train_x, train_y)
```

Best hyperparameters: {'classify__C': 10, 'classify__penalty': 'l2', 'select__k': 8}
Best accuracy score: 0.9897078531224872


Logistic regression model

The accuracy for the test set

The Accuracy is : 98.65 %

Classification report:


	precision	recall	f1-score	support
0	0.99	0.99	0.99	1217
1	0.99	0.98	0.98	856
accuracy			0.99	2073
macro avg	0.99	0.99	0.99	2073
weighted avg	0.99	0.99	0.99	2073





Linear SVC

The accuracy of
the linear SVC
model was
97.35% initially.



However, the
cross-validation
showed a lower
accuracy of
92.97%,
revealing the
model's bias.

Linear SVC

Optimizing the model with GridSearchCV

Best hyperparameters: {'C': 1, 'dual': False, 'loss': 'squared_hinge', 'max_iter': 1000, 'penalty': 'l1'}

Best accuracy score: 0.9867594582520532

The Accuracy is : 97.88 %

Classification report:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	1217
1	0.98	0.97	0.97	856
accuracy			0.98	2073
macro avg	0.98	0.98	0.98	2073
weighted avg	0.98	0.98	0.98	2073

Random Forest

- Set the initial parameter:


```
RandomForestClassifier(n_estimators = 5,max_leaf_nodes =  
5,random_state=42)
```

- The Accuracy is : 99.13 %

Random Forest

```
▶ Ran  
▶ estimator:  
  ▶ Random
```

Best hyperpar
Best accuracy



2', 'n_estimators': 142}

Neural Network

- Best Parameters: {'solver': 'adam', 'max_iter': 1000, 'learning_rate_init': 0.001, 'hidden_layer_sizes': (10, 5), 'alpha': 0.01, 'activation': 'logistic'}

The Accuracy is : 92.23 %

Classification report:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	1217
1	1.00	0.81	0.90	856
accuracy			0.92	2073
macro avg	0.94	0.91	0.92	2073
weighted avg	0.93	0.92	0.92	2073

The Accuracy is : 97.06 %

Classification report:

	precision	recall	f1-score	support
0	0.95	1.00	0.98	1217
1	1.00	0.93	0.96	856
accuracy			0.97	2073
macro avg	0.98	0.96	0.97	2073
weighted avg	0.97	0.97	0.97	2073



Naïve Bayes

- The Naïve Bayes model initially achieved a high accuracy score of 95.8%.
- The model's performance worsened after using GridSearchCV to optimize the var_smoothing parameter.
- Combining the SelectKBest and GridSearchCV significantly improved the accuracy score of 96.33%.



The time train the optimized model

- Logistic regression: 0.05 seconds
- SVM: 0.83 seconds
- Random forest: 0.73 seconds
- Neural network: 1.43 seconds
- Naive Bayes: 0.01 seconds



Conclusion

- The Random Forest model demonstrated the highest accuracy score and f1-scores for both classes, making it the top performer.
- Meanwhile, the naive Bayes is the fastest model to train and predict.
- However, the Logistic Regression model is a suitable choice for predicting and evaluating our dataset in this project.



Thank you

